

MNIST Dataset

Classification, Logistic Regression

Last few weeks

Aurélien Géron, ***Hands-on-Machine Learning***

1. Look at the big picture
2. Get the data and set aside a test set
3. Discover and visualise the data to gain insights
4. Prepare the data for Machine Learning algorithms
5. Identify a suitable metric for evaluating the task
6. Select a model and train it
7. Fine-tune your model
8. Present your solution -> **Final Assignment!**
9. Launch, monitor and maintain your system -> **if ready!**

This week

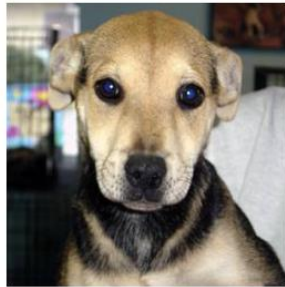
- Classification
- Logistic Regression
- Performance Metrics for classification
- Classification using MNIST dataset (hand-drawn digits)
- Sklearn API

Classification deals with categorical data

- We're used to applying **regression models** to **continuous data** to make predictions (e.g. house prices could be from 100k to 10 mil.).
- But **classification** lends itself to **categorical data** – to classify as a particular category (a value that is one of : yes/no, true/false, north/south, A/B/C/D, 0/1/2/3/4 etc)



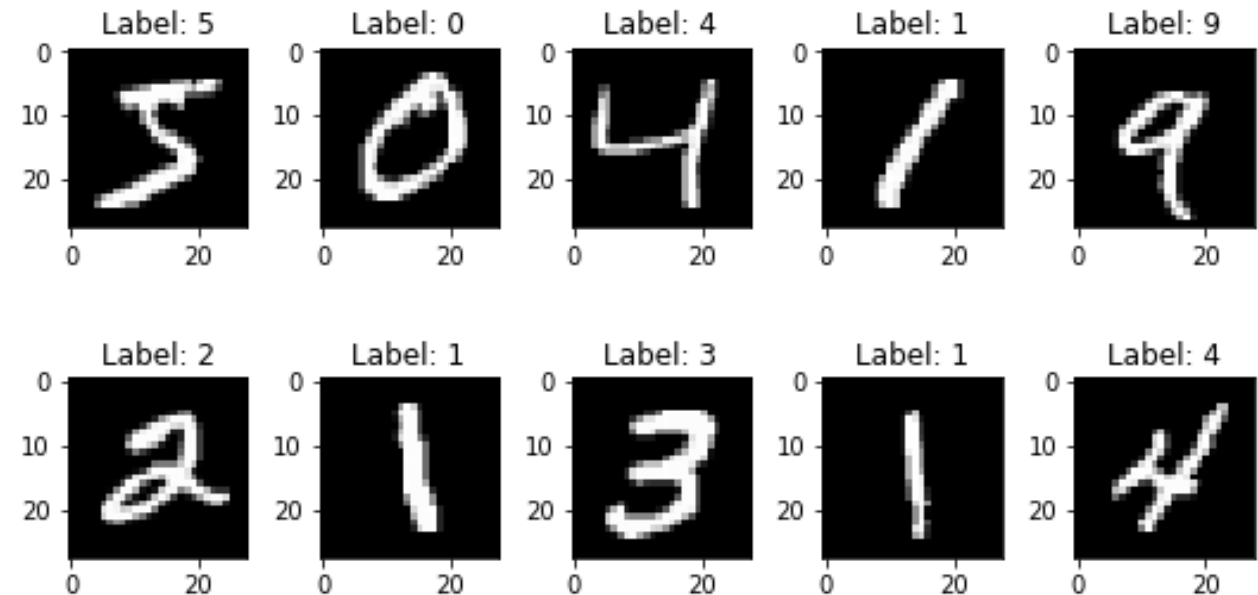
Cat



Dog

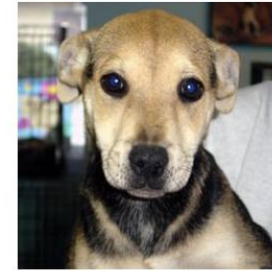


Cat



Classification tasks

- **Binary classification:** two classes (0 or 1)

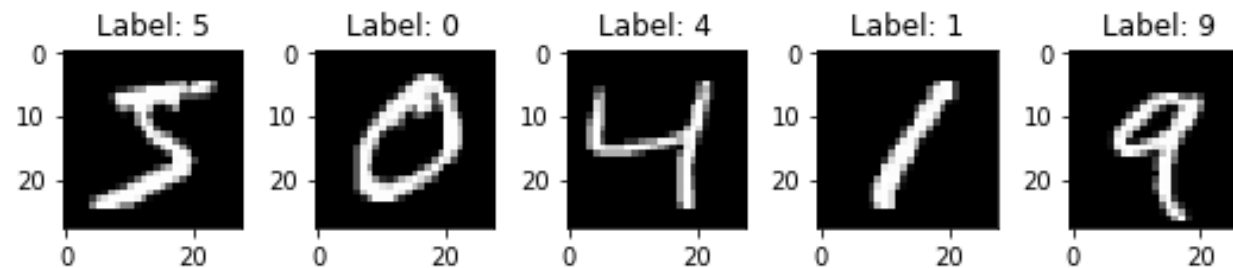


Dog



Cat

- **Multi-class:** (exclusive: can only be one: this is a '9')

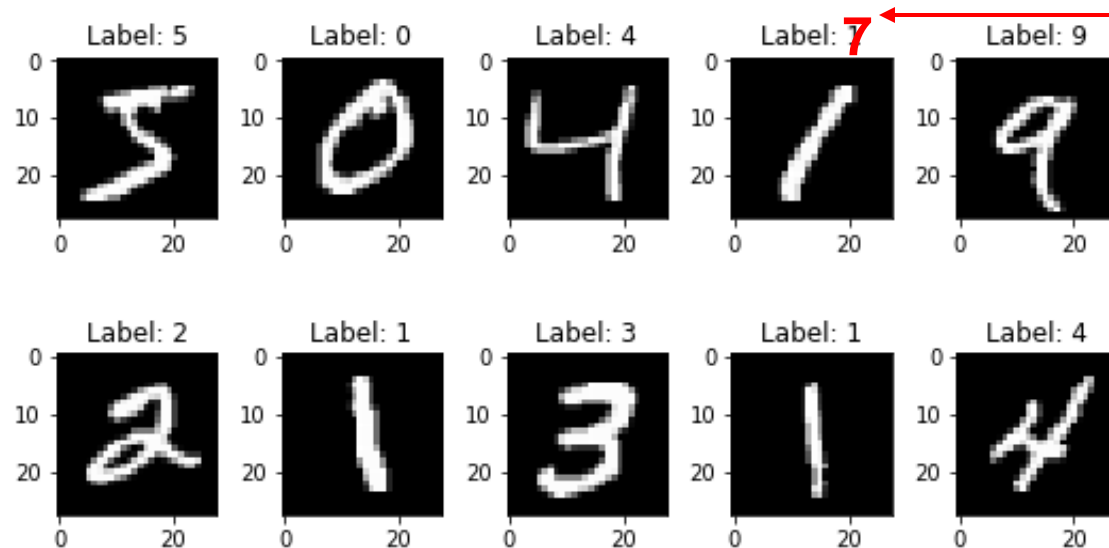


- **Multi-label:** (can belong to more than one label: Action AND Sci-fi)

5. Metrics: Accuracy

- Performance metrics are trickier for classification.
- Accuracy would be the most intuitive obvious one to measure

$$accuracy = \frac{\# \text{ correctly predicted records}}{\# \text{ total records}}$$



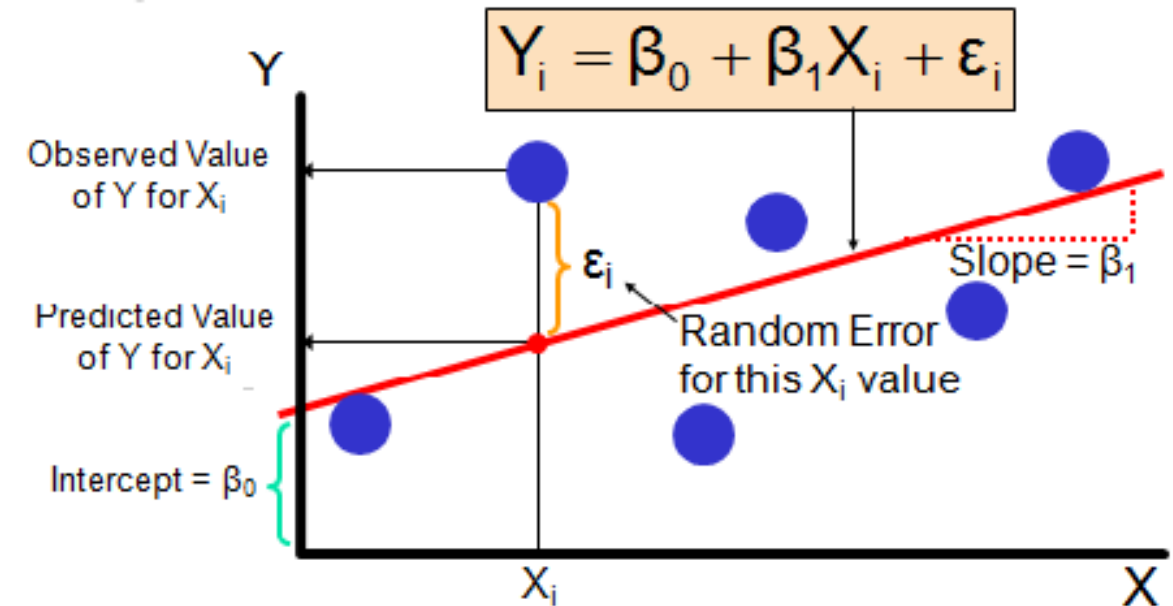
Let's say that the model incorrectly classified this image as a 7 rather than a 1

$$accuracy = \frac{9}{10} = 90\%$$

5. Identify a metric for evaluation

- **Metrics for Regression tasks:**

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- Mean Absolute Error (MAE)
- R^2 (variance explanation)



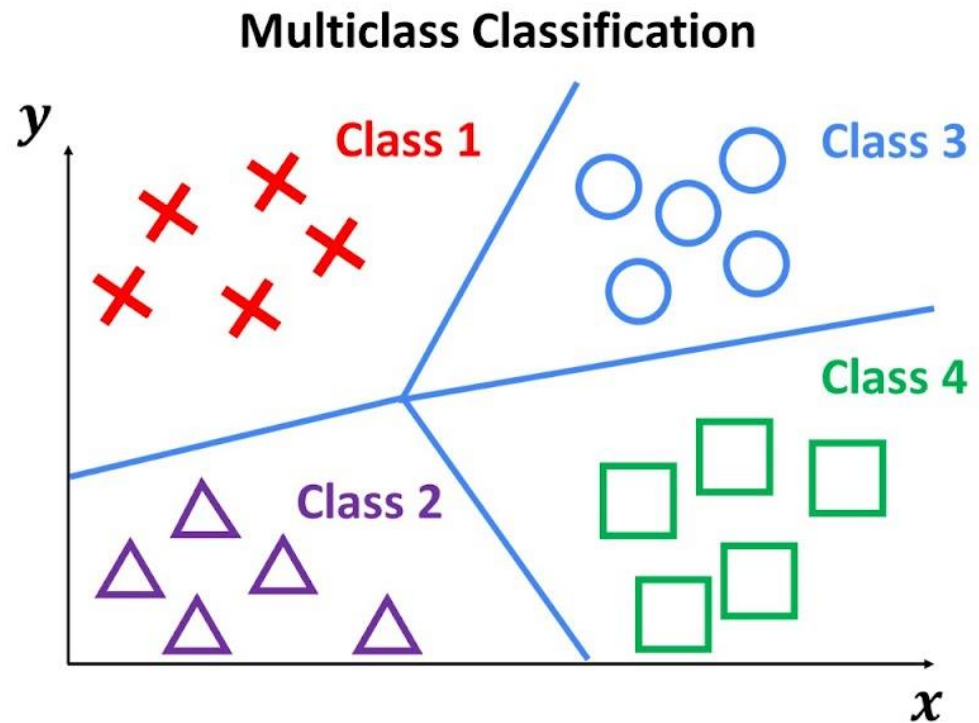
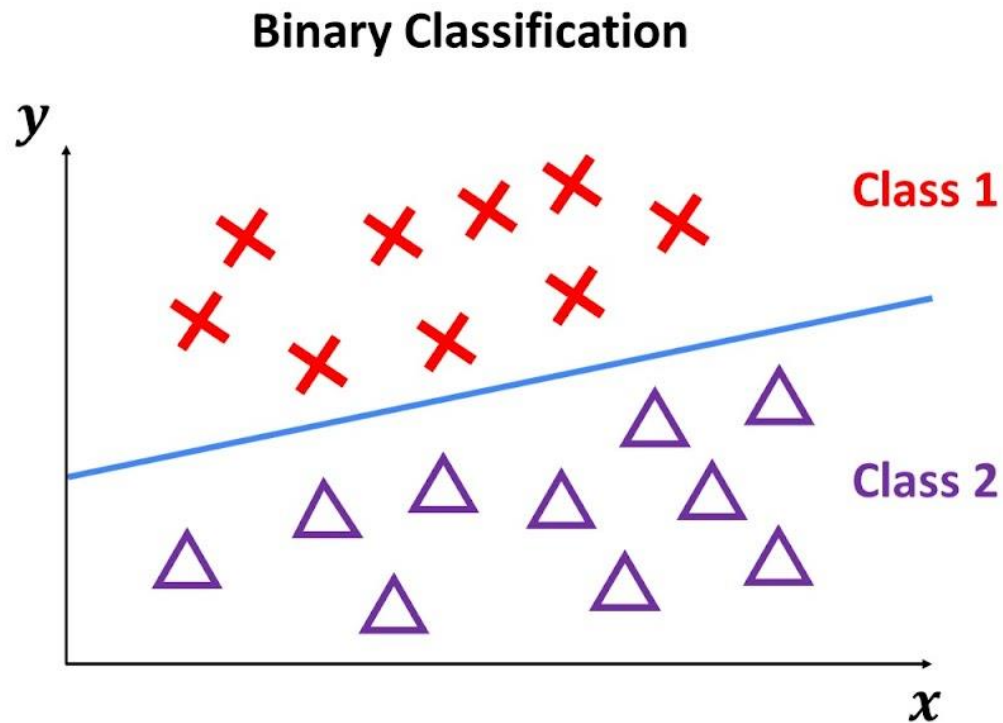
- **Metrics for Classification tasks:**

- Precision → F1-score
- Recall
- **Accuracy** (percentage correct)
- Confusion Matrix (type I and II errors)

		Predicted	
		Spam	Non-spam
Actual	Spam	600 (True positive)	300 (False negative)
	Non-spam	100 (False positive)	9000 (True negative)

Logistic Regression

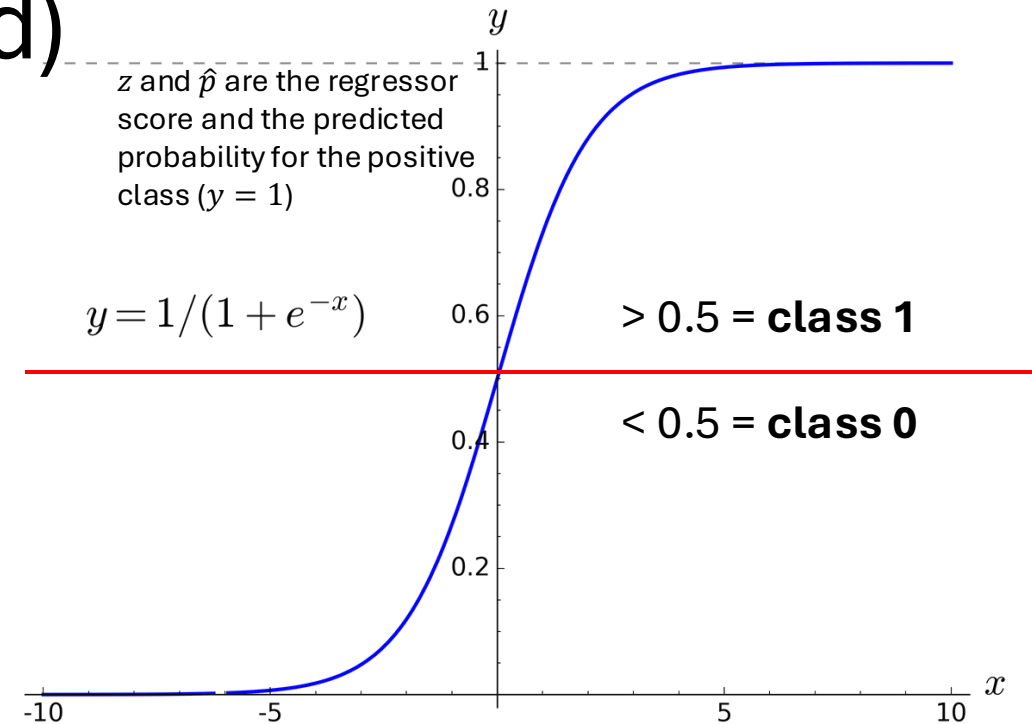
- In classification we would typically use:
 - **sigmoid** (for binary/multi-label)
 - **softmax** (for multi-class) to output probabilities.



$$\hat{p}(\boldsymbol{\theta}) = \hat{p}(\mathbf{w}, b) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-(\mathbf{x}^T \mathbf{w} + b)}}$$

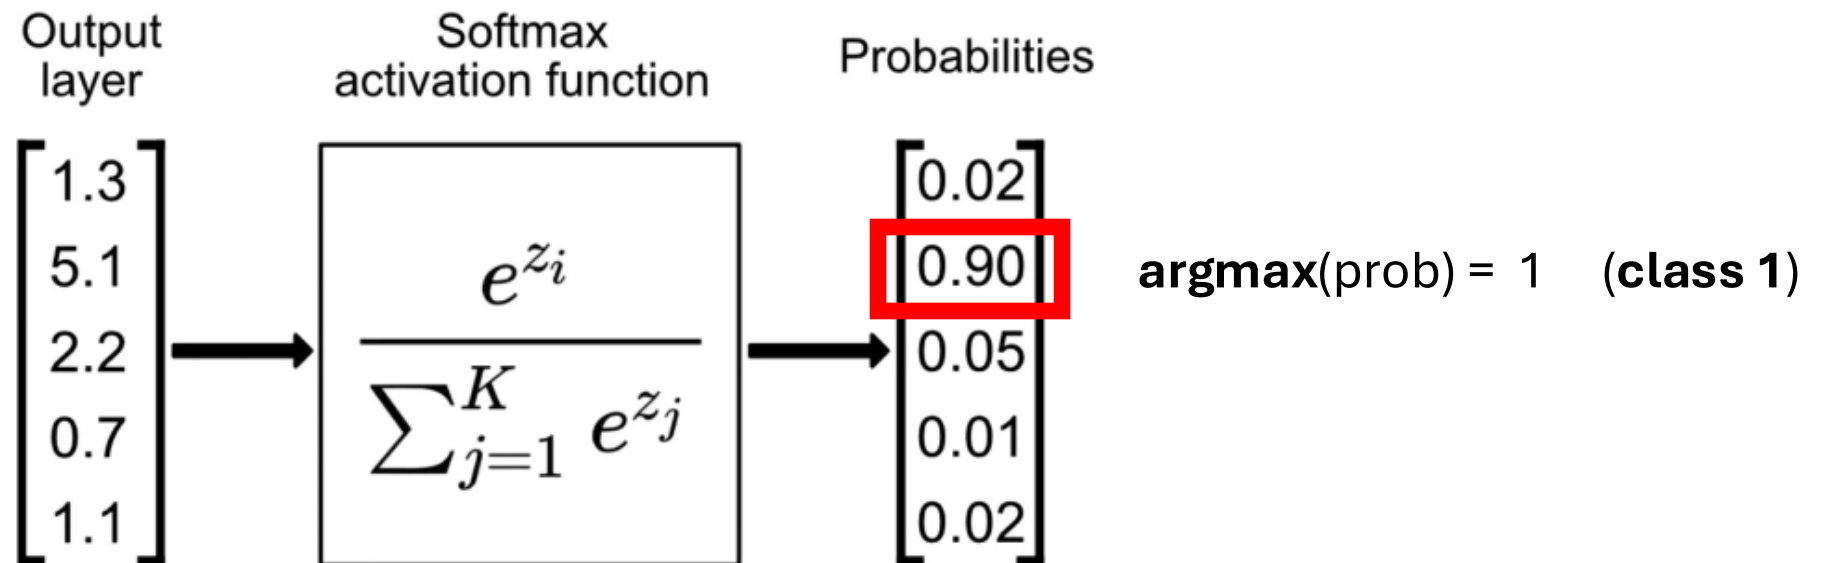
Logistic Regression (Sigmoid)

- Binary classification (class 0 or 1)
- Estimates the probability that a sample belongs to a certain class by training a (linear) regressor that will return scores in the $(-\infty, +\infty)$ interval (continuous)
- Then passing the output of the regressor to a **logistic (sigmoid) function** so that the output will be a value **between 0 and 1**.
- If the **output is > 0.5** assign to **class 1**
- If the **output is < 0.5** , otherwise assign to **class 0**



Logistic Regression (Softmax)

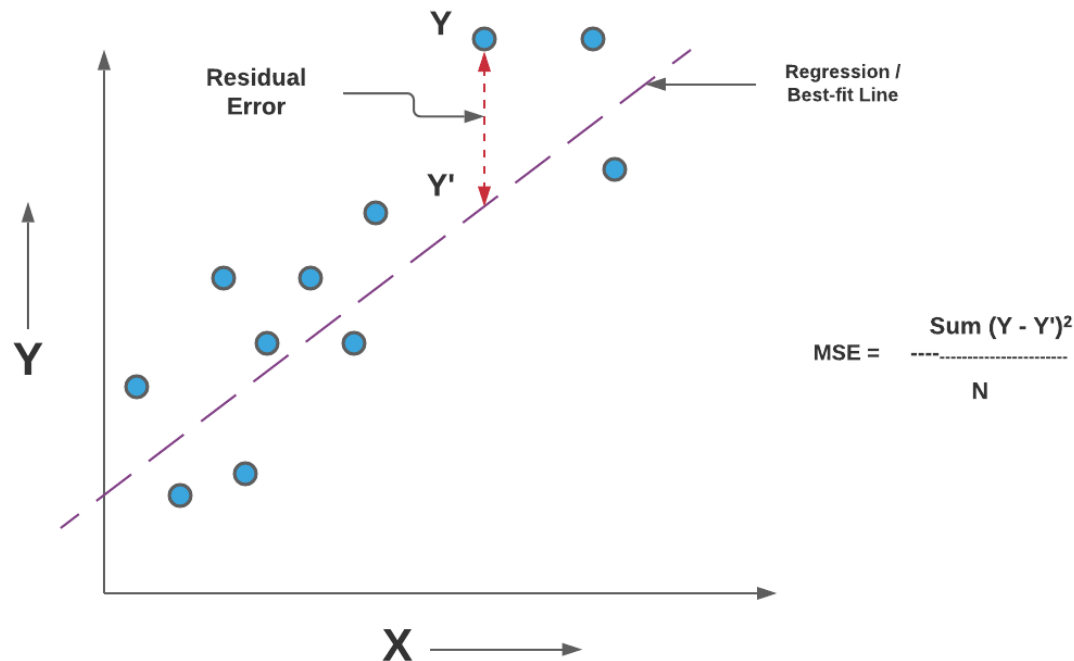
- Multi-class classification (classes 0, 1, 2, 3, 4 etc.)
- **The softmax function** is used in multi-class classification to convert raw scores (logits) into **probabilities** that sum to **1**.
- Then use **argmax()** to find class with highest probability score



5. Metrics: Mean Squared Error (MSE)

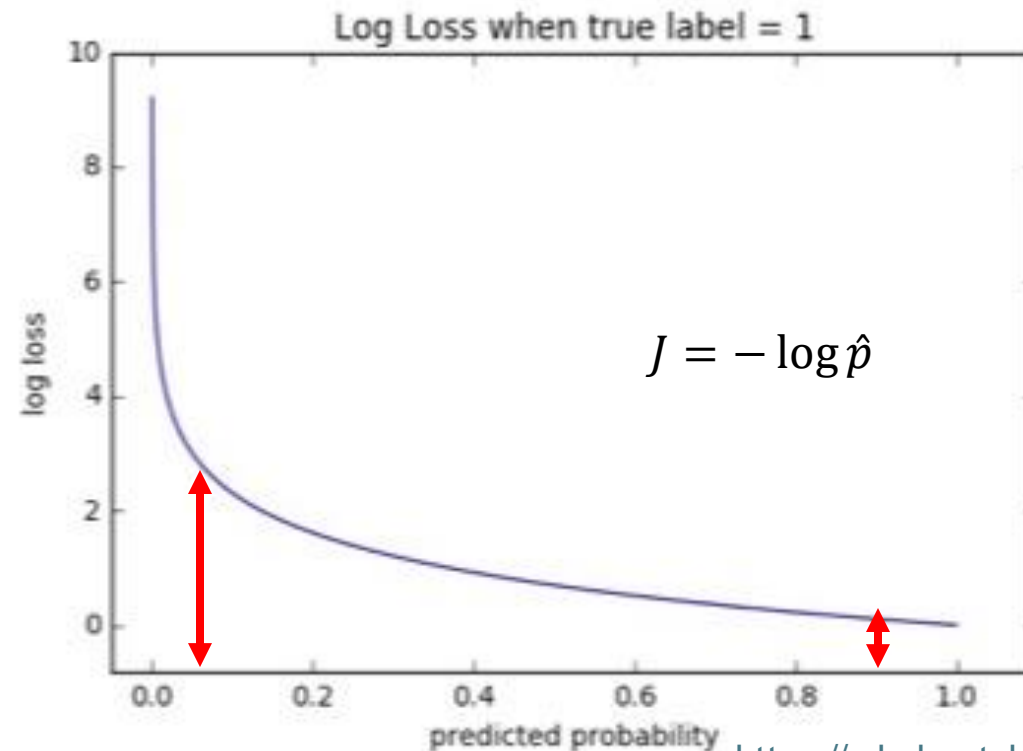
- This is the mean of the squared errors.
- Larger errors are noted more than with MAE, making MSE more popular.

$$\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$



Log loss cost function

- Log loss creates a **logarithmic gradient**, meaning the **penalty for incorrect predictions increases sharply when the model is wrong** and decreases when it is closer to the correct probability.



$$\text{crossentropy} = J(\boldsymbol{\theta}) = -\frac{1}{m} \sum_{i=1}^m \sum_k^K y_k \log \hat{p}_k$$

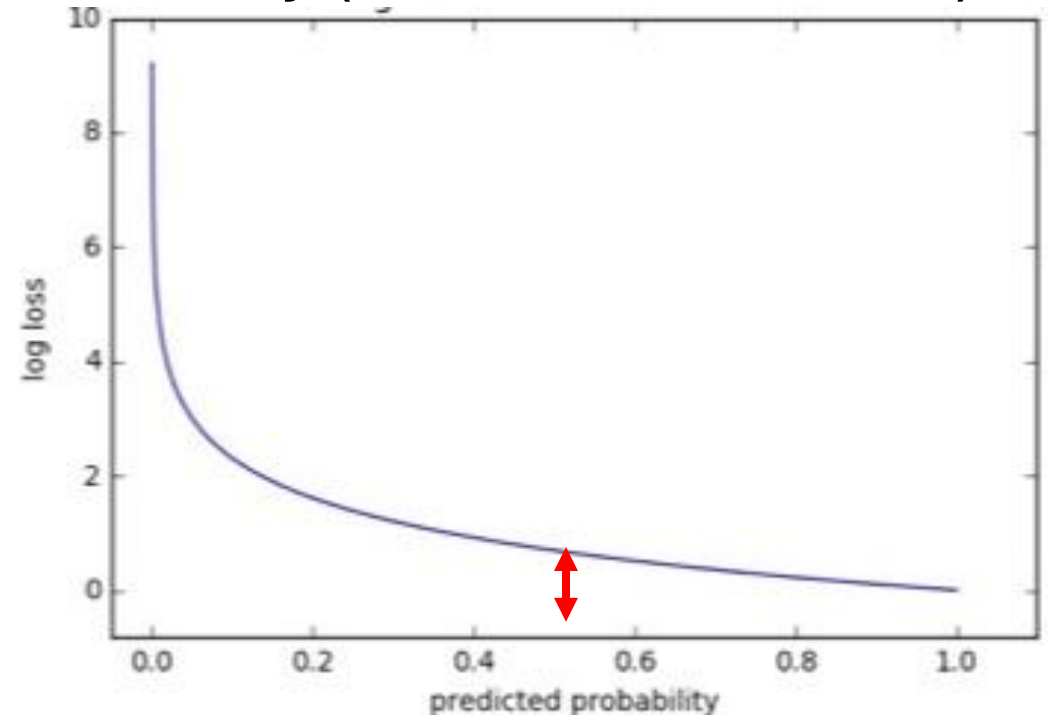
Softmax regression: cross entropy cost function

- The cost function used for softmax regression is a multi-class extension of the **log-loss** function, called **cross-entropy**.
- It applies **log-loss** to the highest probability (and selected class).

Softmax()

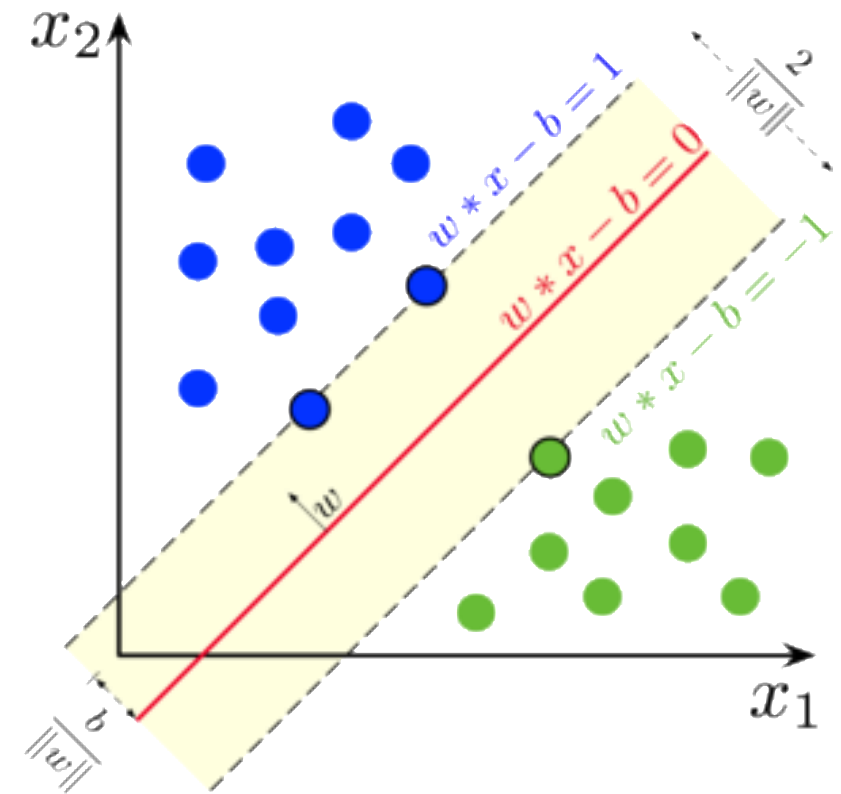
0.2
0.3
0.5

$$L = -\log(\mathbf{0.5}) = 0.693$$



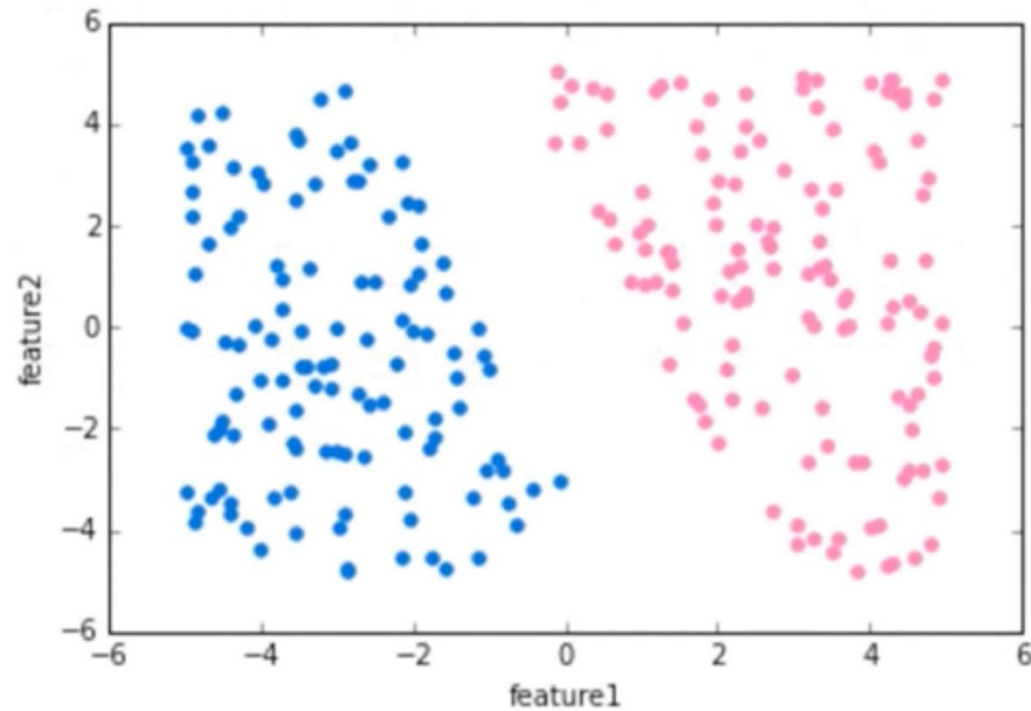
Support Vector Machines (SVMs)

- An SVM constructs a hyper-plane (or set of hyper-planes) in a high dimensional space, which can be used for classification (or regression or other tasks).
- Intuitively, a good separation is achieved by the hyper-plane that has the **largest distance** to the **nearest training data points** of any class (so-called **functional margin**),
- In general, the **larger the margin**, the **lower the generalization error of the classifier**
- Let's look at an example...



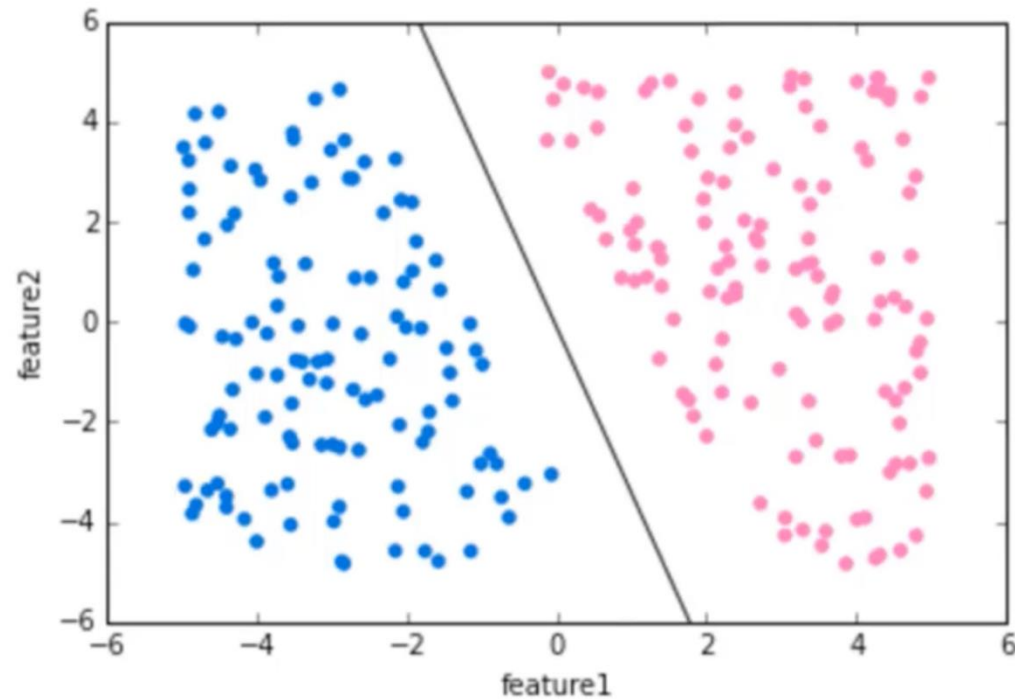
SVM Example:

- Here are two dimensions, and we have two classes / labels



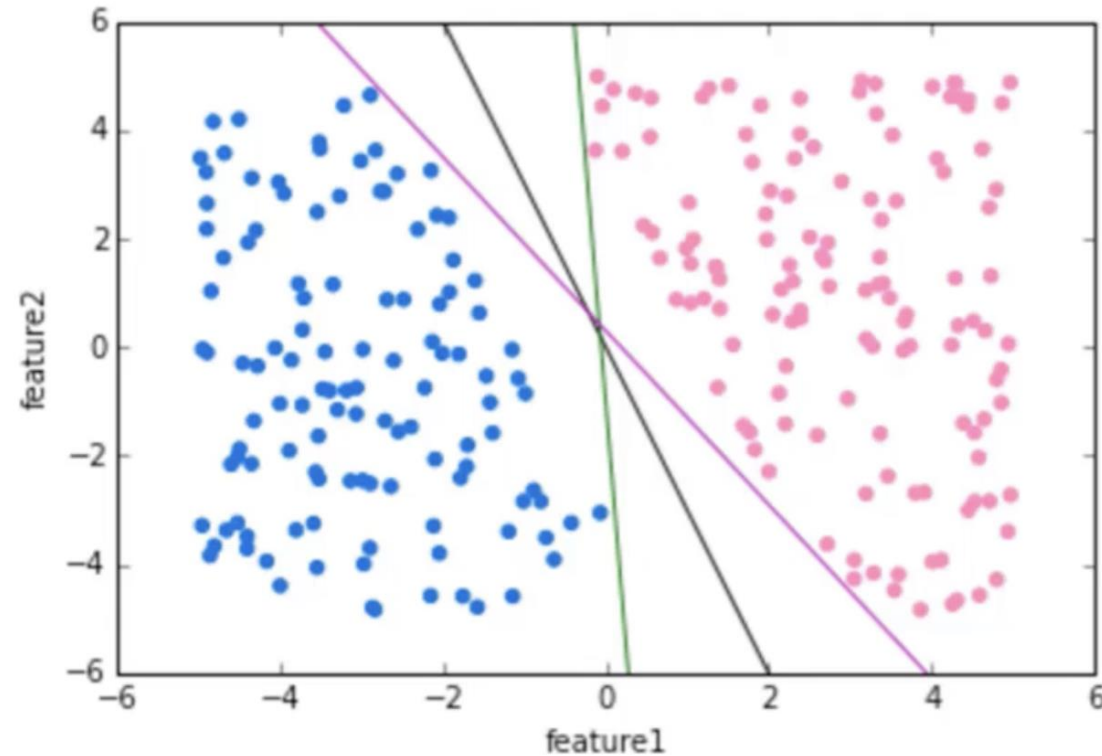
SVM Example:

- We could draw a regression dividing line to help us classify



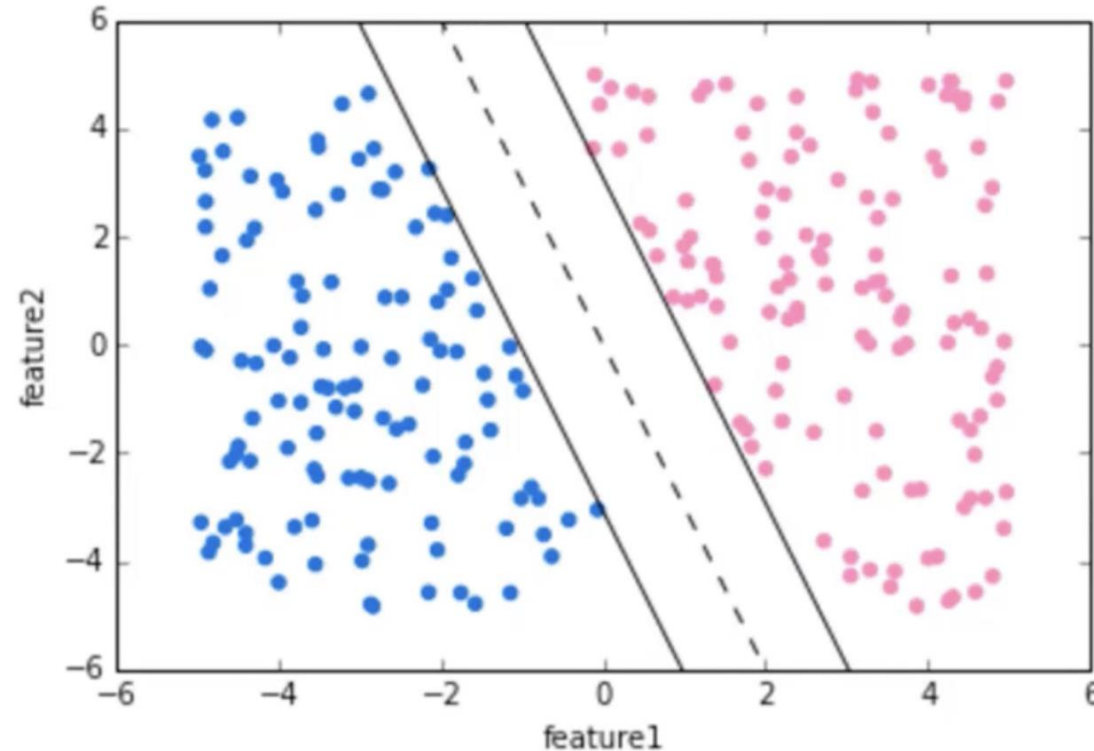
But there are many options...

- Which slope/gradient should we go with?



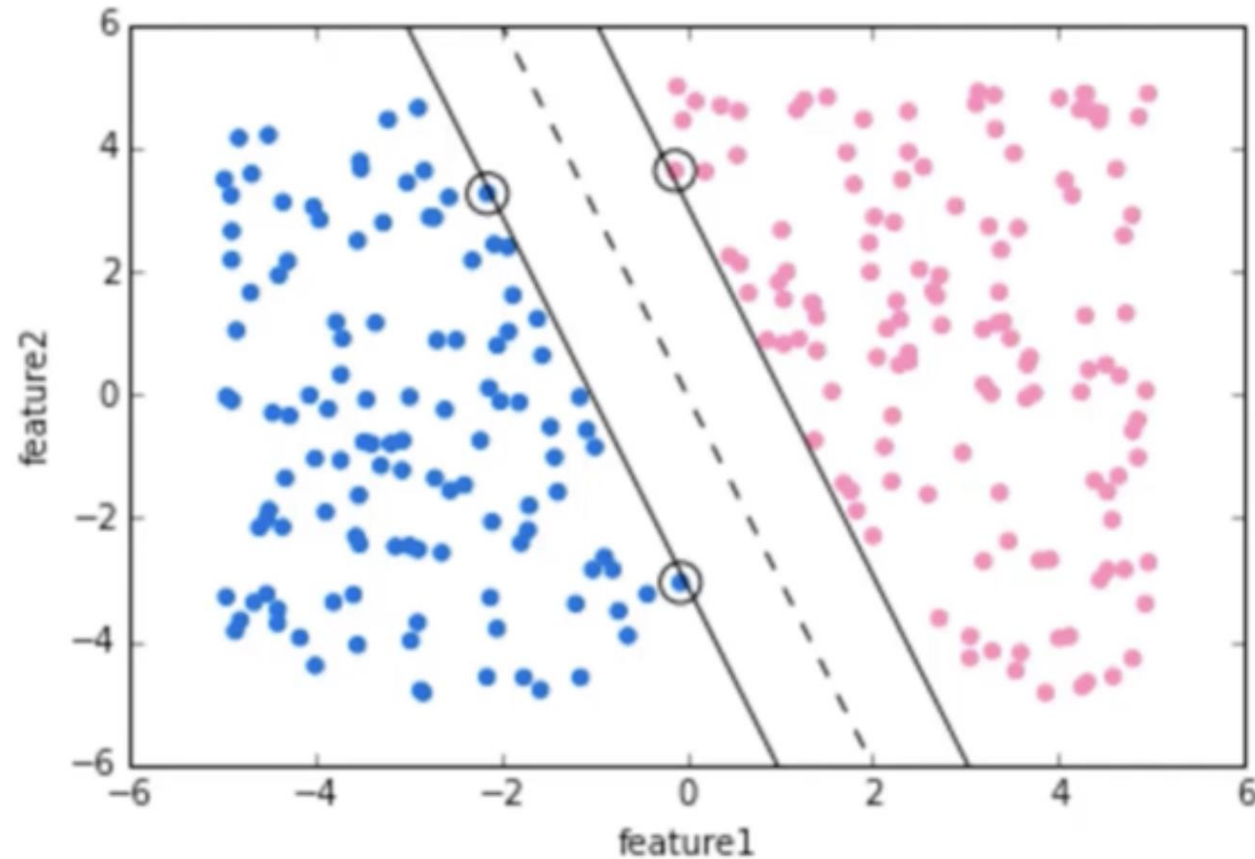
SVM maximises margin between classes

- We would like to choose a hyperplane that maximises the margin between classes – reduces the likelihood of error



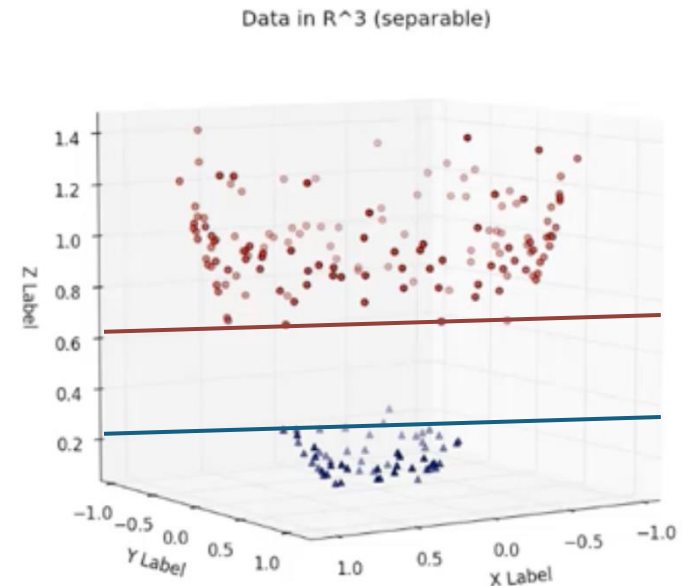
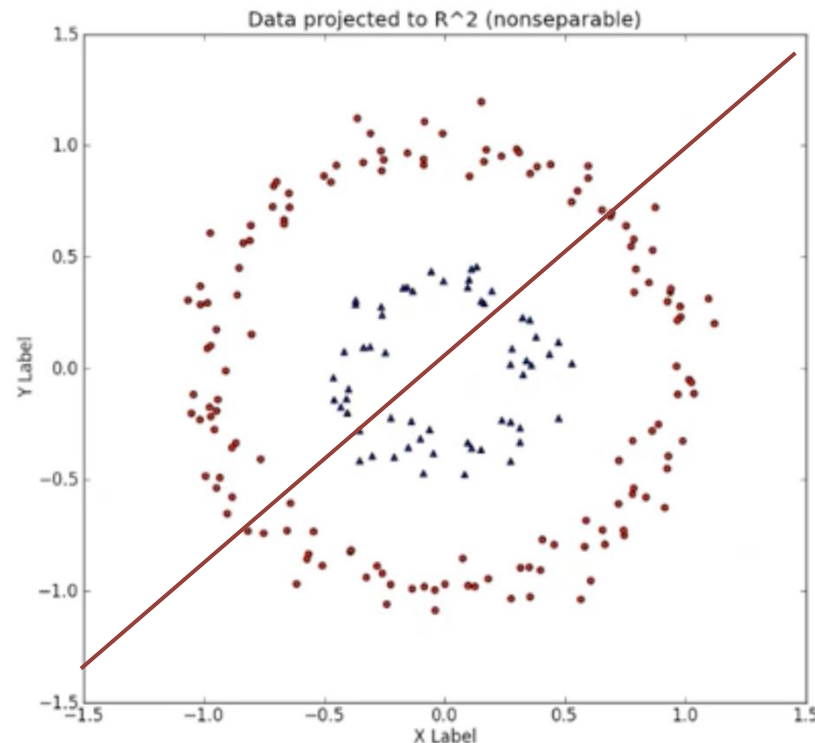
Support Vectors

- **Points** that support the margin lines are known as **Support Vectors**



Support Vectors: the 'Kernel Trick'

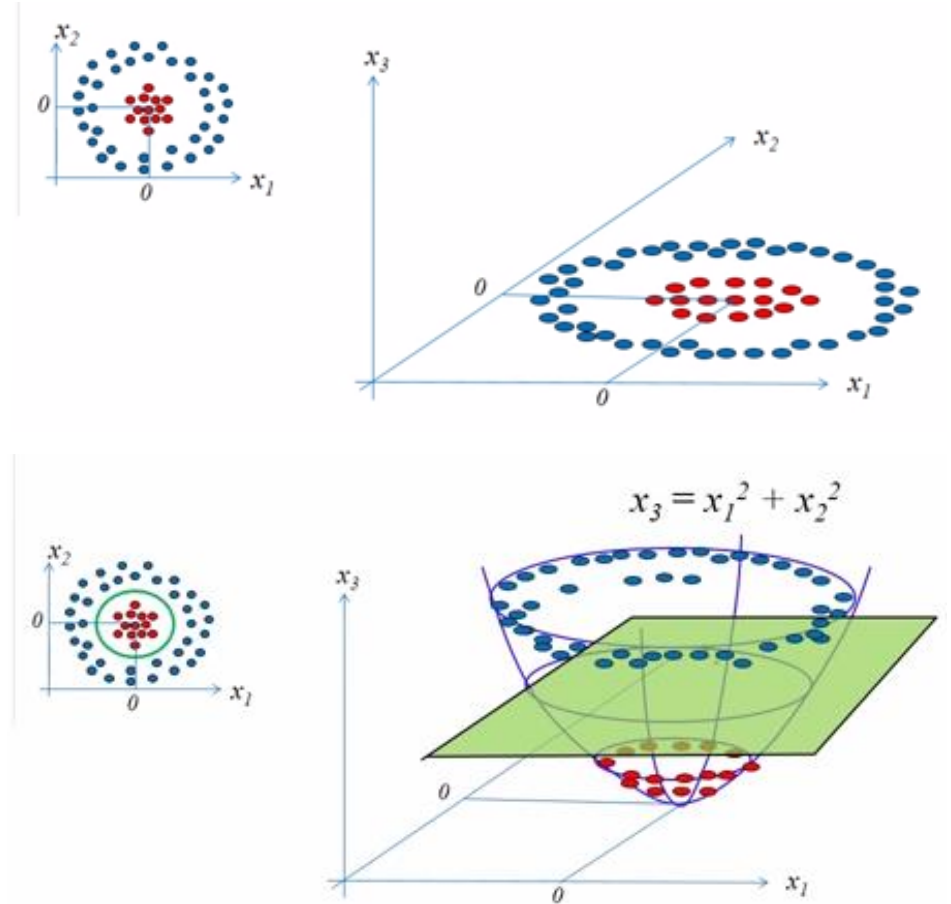
- This concept can expand to **non-linear data** (that isn't separable with a straight line) through a 'kernel-trick'
- In 2D, a straight cannot be drawn to separate the classes, but in 3D, it can!!



Support Vectors: the ‘Kernel Trick’

- To solve a nonlinear problem with SVM:
 1. We **transform** the training data onto a **higher dimensional feature space** via a **mapping function ϕ** .
 2. We **train a linear SVM model** to classify the data in **this new feature space**.
 3. Then, we can **use the same mapping function ϕ** to **transform unseen data** to classify it using the linear SVM model.

The **kernel trick** avoids the explicit mapping that is needed to get linear learning algorithms to learn a nonlinear function or decision boundary.



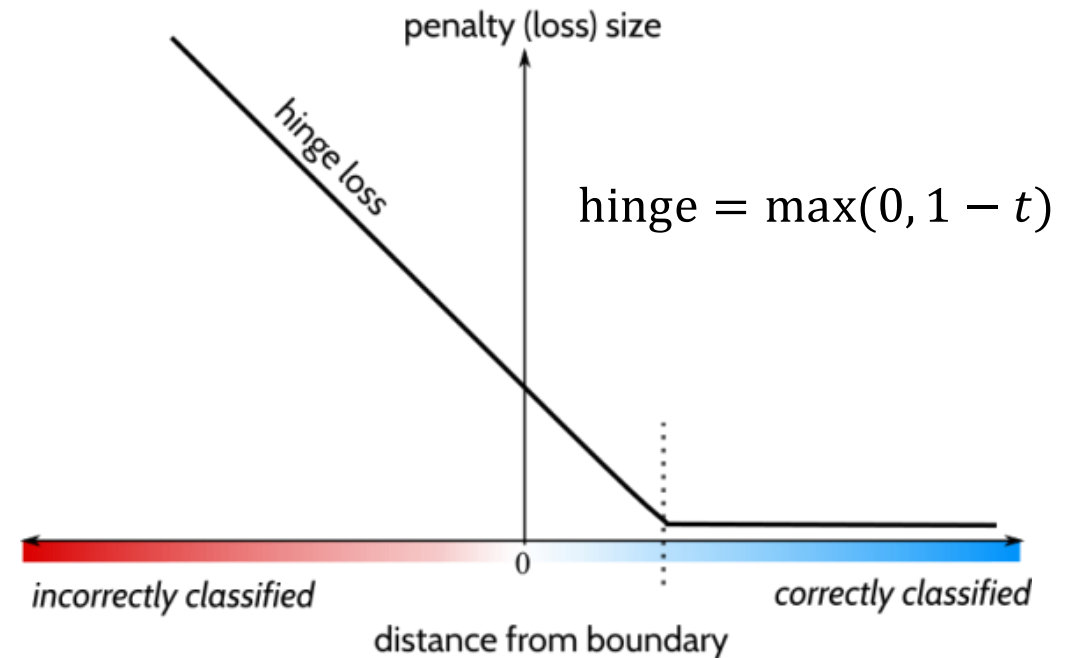
Online SVMs with the hinge loss

$$J(\theta) = J(\mathbf{w}, b) = \underbrace{\frac{1}{2} \mathbf{w}^T \mathbf{w}}_{\text{regularization term}} + C \underbrace{\sum_{i=1}^m \max(0, 1 - t_i(\mathbf{w}^T \mathbf{x}_i + b))}_{\text{hinge loss term}}$$

Traditional SVM are trained offline (batch-training)

However, we can train online SVMs using gradient descent, just like we train logistic or softmax regression classifiers

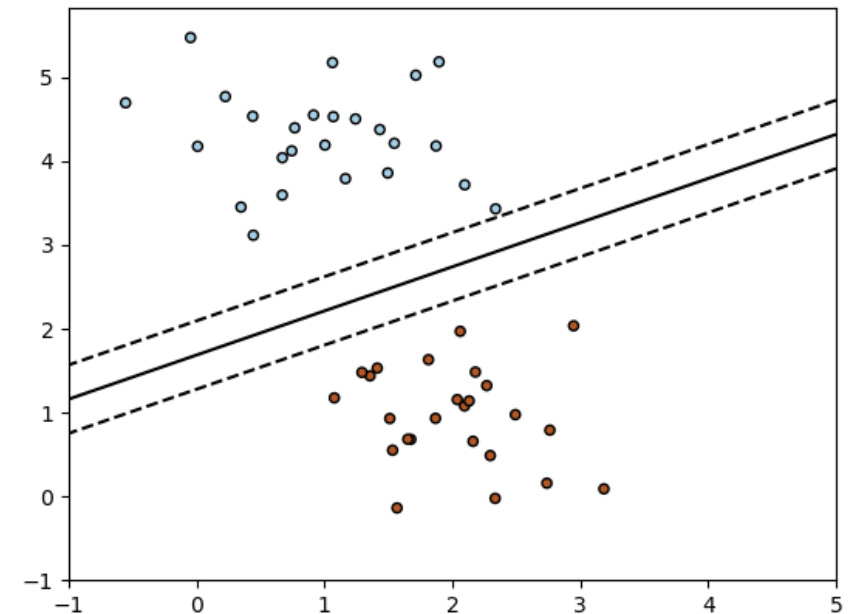
Rather than using the **log loss** we use the **hinge function** as our cost function



Stochastic Gradient Descent Classification

- The class **SGDClassifier** implements a plain **stochastic gradient descent learning** routine which supports different loss functions and penalties for classification.
- The default function scikit-learn is the **Linear SVM decision function**

<https://scikit-learn.org/stable/modules/sgd.html>



https://scikit-learn.org/stable/auto_examples/linear_model/plot_sgd_separating_hyperplane.html

Confusion Matrix

- **Accuracy** = Total correctly classified (True Positive and True Negative) out of all predictions (total population)

- **Precision** = Measure of **predictive positive cases** (True Positive and False Positive). Precision focusses on quality

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

- **Recall (Sensitivity)** = Measure of **actual positive cases** (True Positive and False Negative)

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

- **F1 Score** = Harmonic mean of precision and recall. Provides a balanced score.

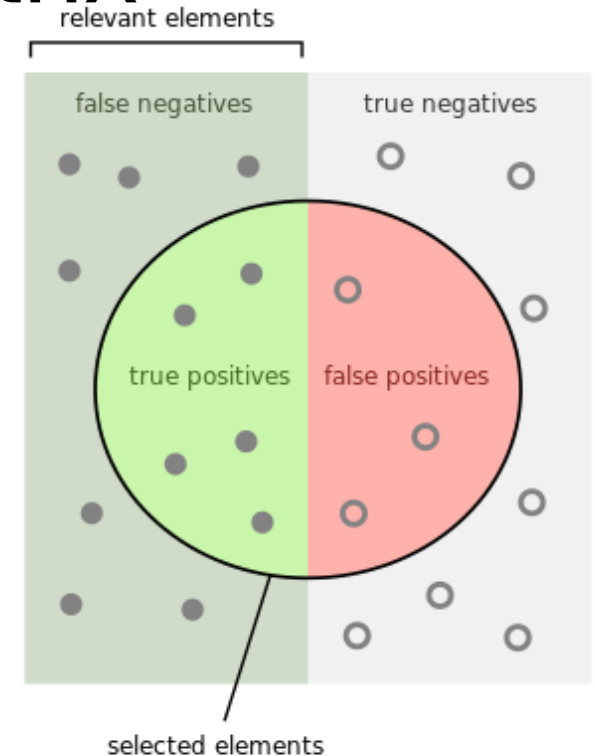
Performance Metrics: Confusion Matrix

$$\text{accuracy} = \frac{\text{correct predictions}}{\text{total predictions}} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{recall} = \frac{TP}{TP + FN}$$

$$F1 = 2 \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$



How many selected items are relevant?



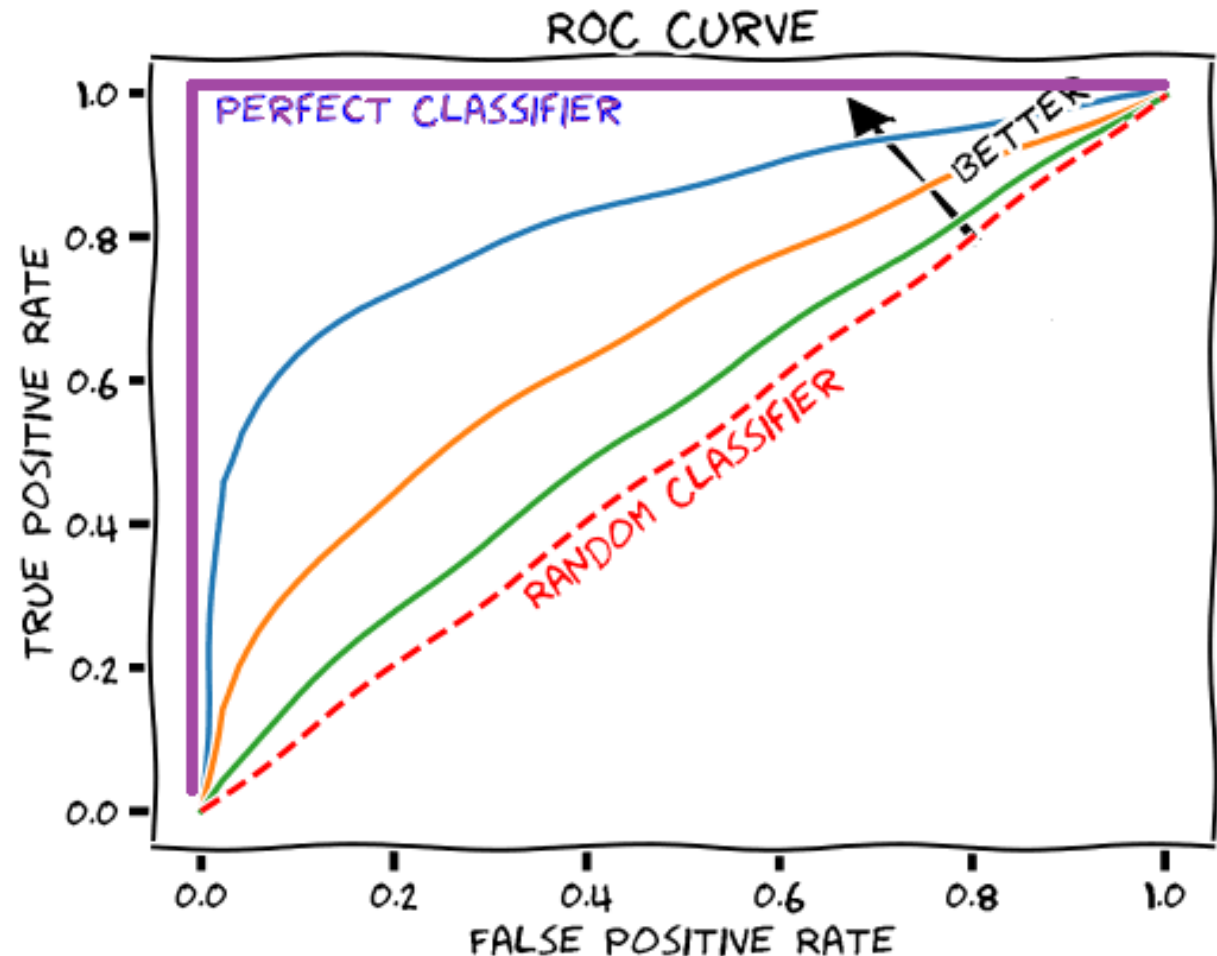
How many relevant items are selected?



Performance Metrics: Area under the ROC curve

$$TPR = recall = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{FP + TN}$$



MNIST Dataset
Sklearn

Images are Numbers



```
[ [ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  3. 18. 18. 18. 126. 136. 175. 26.]  
[ 94. 154. 170. 253. 253. 253. 253. 253. 225. 172.]  
[253. 253. 253. 253. 253. 253. 253. 251. 93. 82.]  
[253. 253. 253. 253. 198. 182. 247. 241. 0. 0.]  
[107. 253. 253. 205. 11. 0. 43. 154. 0. 0.]  
[ 1. 154. 253. 90. 0. 0. 0. 0. 0. 0.]  
[ 0. 139. 253. 190. 2. 0. 0. 0. 0. 0.]  
[ 0. 11. 190. 253. 70. 0. 0. 0. 0. 0.]  
[ 0. 0. 35. 241. 225. 160. 108. 1. 0. 0.]  
[ 0. 0. 0. 81. 240. 253. 253. 119. 25. 0.]  
[ 0. 0. 0. 0. 45. 186. 253. 253. 150. 27.]  
[ 0. 0. 0. 0. 0. 16. 93. 252. 253. 187.]  
[ 0. 0. 0. 0. 0. 0. 0. 249. 253. 249.]  
[ 0. 0. 0. 0. 46. 130. 183. 253. 253. 207.]  
[ 0. 0. 39. 148. 229. 253. 253. 253. 250. 182.]  
[ 24. 114. 221. 253. 253. 253. 253. 201. 78. 0.]  
[213. 253. 253. 253. 253. 198. 81. 2. 0. 0.]  
[253. 253. 253. 195. 80. 9. 0. 0. 0. 0.]  
[253. 244. 133. 11. 0. 0. 0. 0. 0. 0.]  
[132. 16. 0. 0. 0. 0. 0. 0. 0. 0.]  
[ 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]  
[ 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]  
[ 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

- Here is a hand drawn digit 5 from the MNIST library [28 x 28] px
- An image is just a matrix of numbers [0,255] 0 for white, 255 for black

Logistic Regression with sklearn

- from **sklearn.model_selection** import **train_test_split**
- from **sklearn.linear_model** import **LogisticRegression**
- from **sklearn.metrics** import **confusion_matrix**
- from **sklearn.metrics** import **classification_report**

Logistic Regression – fit the model to the data

```
from sklearn.linear_model import LogisticRegression
```

```
log_model = LogisticRegression()
```

```
log_model.fit(X_train, y_train)
```

```
log_model.classes_
```

```
array([0, 1])
```

Logistic Regression – let's test the model

```
predictions = log_model.predict(X_test)
```

```
array([0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0,  
       0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,  
       0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0,  
       0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1,  
       1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0,  
       0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
       0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,  
       0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1,  
       1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1,  
       0, 0])
```

Logistic Regression – let's evaluate the model

from **sklearn.metrics** import **confusion_matrix**

confusion_matrix(y_test, predictions)

```
array([[90, 19],  
       [48, 43]])
```

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

Logistic Regression – let's evaluate the model

```
from sklearn.metrics import classification_report
```

```
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.69	0.92	0.79	101
1	0.88	0.58	0.70	99
accuracy			0.75	200
macro avg	0.78	0.75	0.74	200
weighted avg	0.78	0.75	0.74	200

Coming up: Decision Trees

- Decision Trees (DTs) are a non-parametric supervised learning method used for classification (and regression).
- The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features.

Decision tree trained on all the iris features



Coming up: Neural Networks for classification

